

In [1]:

```
import numpy as np
import pandas as pd
import pickle
from Bio import SeqIO
import ipdb
import random
random.seed(10)
import re

class DataFormat():
    '''Class for merging all of the data from the phagescope database

    ***Put all files into one folder with no changes in the file names***

    ***file Names:

    -TemPhD # folder with all of the fasta files
    -TemPhD_antimicrobial_resistance_gene_data.tsv
    -temphd_phage_annotated_protein_meta_data.tsv
    -temphd_phage_meta_data.tsv
    -TemPhD.gff3

    ****'''

    def __init__(self,folderpath, DBName):
        '''Initialization takes the name of the database data and the pth to the folder

        self.path = folderpath #Path to data folder
        self.DB = DBName #Name of database

        self.amr_tsv = '%s_antimicrobial_resistance_gene_data' % self.DB
        self.protein_tsv = '%s_phage_annotated_protein_meta_data' % self.DB.lower()
        self.gff3_file = '%s' % self.DB
        self.meta_tsv = '%s_phage_meta_data'% self.DB.lower()

        self.data_df = None #Data frame with all of the merged data
        self.fasta_path = '%s/%s' % (self.path,self.DB)#fasta folder path

    #Functions for loading and writing files
    def __load_tsv__(self,file_name):
        '''Function for loading .tsv files'''

        path = '%s/%s.tsv' % (self.path,file_name)
        df = pd.read_csv(path,sep="\t",dtype=str)
        return df

    def __load_gff3__(self,file_name):
        '''Function for loading .gff3 file'''

        path = '%s/%s.gff3' % (self.path,file_name)
        df = pd.read_csv(path,sep="\t",header=None,comment="#",names=["seqid", "source",
                           "score", "strand", "phase", "attributes"])
        return df
```

```

def __load_fasta__(self,file_name):
    '''Function for loading .fasta file'''

    path = '%s/%s.fasta' % (self.fasta_path,file_name)
    fast_list = list(SeqIO.parse(path, "fasta"))

    df = pd.DataFrame([{"id": rec.id,"description": rec.description,"sequence": str(rec.seq),"length": len(rec.seq)} for rec in fast_list])
    return df

def load_data_pickle(self, file_name):
    '''Load Pickle File'''

    path = '%s/%s.pickle' % (self.path,file_name)
    f = open(path, 'rb')
    data = pickle.load(f)
    f.close()

    return data

def savePickle(self,file_name,data):
    '''Save Data as Pickle'''

    path = '%s/%s.pickle' % (self.path,file_name)
    f = open(path, 'wb')
    pickle.dump(data, f)
    f.close()

    return

def export_csv(self,file_name,data):
    '''export data as .csv'''

    path = '%s/%s.csv' % (self.path,file_name)
    data.to_csv(path)

    return

```

#Tools

```

def reverse_complement(self,seq):
    """Return the reverse complement of a DNA sequence."""
    complement = str.maketrans("ACGTacgt", "TGCATgca")
    return seq.translate(complement)[::-1]

def __get_Phage_seqs__(self,phage_ids,protein_df):
    '''returns all phage sequences and protein sequences from list of phage_ids'''


    df_seqs = self.__load_fasta__(phage_ids[0])
    for x in range(1, len(phage_ids)):
        new_row = self.__load_fasta__(phage_ids[x])
        df_seqs = pd.concat([df_seqs, new_row], ignore_index=True)

    # build a lookup dict
    phage_to_seq = dict(zip(df_seqs["id"], df_seqs["sequence"]))

```

```

phage_to_desc = dict(zip(df_seqs["id"], df_seqs["description"]))
phage_to_len = dict(zip(df_seqs["id"], df_seqs["length"]))

rows = []

for x, pro_row in protein_df.iterrows():
    phage_id = pro_row["Phage_ID"]
    start = int(pro_row["Start"])
    stop = int(pro_row["Stop"])
    strand = pro_row["Strand"]
    prot_id = pro_row.get("Protein_ID")

    genome_seq = phage_to_seq.get(phage_id)
    if genome_seq is None:
        dna_seq = None
    else:
        # Prodigal/GFF: [start, stop] is 1-based inclusive
        start_i = start - 1 # 0-based
        stop_i = stop         # exclusive for Python slicing
        dna_seq = genome_seq[start_i:stop_i]
        if strand == "-":
            dna_seq = self.reverse_complement(dna_seq)

    rows.append(
    {
        "Protein_ID": prot_id,
        "sequence": dna_seq,
        "description": phage_to_desc.get(phage_id),
        "length": phage_to_len.get(phage_id),
    }
)
df_out = pd.DataFrame(rows)

return df_out

def __get_phages_with_no_AMR__(self, phage_AMR_ids, meta_data, n):
    '''Generates a list of non-AMR phage ids'''
    non_AMR_ids = []

    meta_data_AMR_removed = meta_data[~meta_data["Phage_ID"].isin(phage_AMR_ids)].reset_index()
    meta_data_AMR_removed_ids = meta_data_AMR_removed["Phage_ID"].tolist()

    non_AMR_ids = random.sample(meta_data_AMR_removed_ids, n)

    return non_AMR_ids

def __get_phages_with_AMR__(self, phage_AMR_ids, meta_data, n):
    '''generate a list of AMR ids'''

    AMR_ids = []

    meta_data_AMR = meta_data[meta_data["Phage_ID"].isin(phage_AMR_ids)].reset_index()
    meta_data_AMR_ids = meta_data_AMR["Phage_ID"].tolist()

    AMR_ids = random.sample(meta_data_AMR_ids, n)

```

```

    return AMR_ids

def __remove_repeated_cols_df__(self,df):
    ''' Drop duplicate columns from data frame merger'''

    unique_cols = []
    seen = set()

    for col in df.columns:
        # Remove typical merge suffixes: _x, _y, _left, _right
        base = re.sub(r'(_x$|_y$|_left$|_right$)', '', col, flags=re.IGNORECASE)
        key = base.lower()

        if key not in seen:
            seen.add(key)
            unique_cols.append(col)

    return df[unique_cols]

def __drop_index_columns__(self,df):
    # drop things like 'index', 'index_x', 'index_y', 'level_0', 'Unnamed: 0', etc.
    bad = [c for c in df.columns
           if c.lower().startswith("index")
           or c.lower().startswith("level_0")
           or c.lower().startswith("unnamed")]
    return df.drop(columns=bad, errors="ignore")

#Functions for creating datafile
def create_data_file(self, number_of_AMR_seqs = None, Number_of_control_seqs = None,
                     high_quality_only = False, high_med_quality_only = True, save_f
                     AMR_proteins_Only = False):

    '''Main function for creating datasets:

    Inputs:
        number_of_AMR_seqs = None # number of phage AMR sequences in the dataset, default is 0
        Number_of_control_seqs = None # number of phage non-AMR sequences, default is 0
        high_quality_only = False #Filter for reads that are high-quality only
        high_med_quality_only = True #filter for reads that are medium and high quality
        save_fname = None #export file name, Also doubles as True for save data, data is saved as pickle
        AMR_proteins_Only = True #only AMR protein sequences, overrides number_of_AMR_seqs

    outputs:
        saves a .pickle file if save_fname is passed
        initializes object with data in self.data_df'''


    #Load in all data files
    phage_AMR_df = self.__load_tsv__(self.amr_tsv)
    gff3_df = self.__load_gff3__(self.gff3_file)
    protein_df = self.__load_tsv__(self.protein_tsv)
    meta_df = self.__load_tsv__(self.meta_tsv)

    #get All AMR phage data
    phage_AMR_ids = phage_AMR_df['Phage_id']

    #filter based on completeness, Can also filter based on score in gff3 but this se

```

```

    if high_quality_only is True:
        meta_df = meta_df[meta_df["Completeness"] == "High-quality"].reset_index()
    elif high_med_quality_only is True:
        meta_df = meta_df[(meta_df["Completeness"] == "High-quality") | (meta_df["Co

#get a number of AMR phage ids
if number_of_AMR_seqs is None:
    number_of_AMR_seqs = len(phage_AMR_ids)
phage_AMR_ids_pos = self.__get_phages_with_AMR__(phage_AMR_ids,meta_df, number_o

#get a number of non AMR phage ids
if Number_of_control_seqs is not None:
    phage_control_ids = self.__get_phages_with_no_AMR__(phage_AMR_ids,meta_df, N
else:
    phage_control_ids = []

#Merge control and positive list
full_list = phage_AMR_ids_pos+phage_control_ids

#filter all data
gff3_df = gff3_df.rename(columns={"seqid":"Phage_ID", "start":"Start","end":"Sto
gff3_df = gff3_df[gff3_df["Phage_ID"].isin(full_list)].reset_index()
protein_df = protein_df[protein_df["Phage_ID"].isin(full_list)].reset_index()
meta_df = meta_df[meta_df["Phage_ID"].isin(full_list)].reset_index()

#Load fasta files
seq_df = self.__get_Phage_seqs__(full_list,protein_df=protein_df)
seq_df = seq_df.rename(columns={"id":"Phage_ID"})

#drop index cols
gff3_df=self.__drop_index_columns__(gff3_df)
protein_df=self.__drop_index_columns__(protein_df)
meta_df=self.__drop_index_columns__(meta_df)
seq_df= self.__drop_index_columns__(seq_df)
#merge dfs
df_one = meta_df.merge(protein_df, on="Phage_ID", how="left")

df_two = seq_df.merge(df_one,on="Protein_ID", how="left")

# numeric coordinates
for col in ["Start", "Stop"]:
    gff3_df[col] = gff3_df[col].astype("int64")
    df_two[col] = pd.to_numeric(df_two[col], errors="coerce").astype("Int64")

# string IDs / strand
for col in ["Phage_ID", "Strand"]:
    gff3_df[col] = gff3_df[col].astype(str).str.strip()
    df_two[col] = df_two[col].astype(str).str.strip()

merge_keys = ["Phage_ID", "Start", "Stop", "Strand"]

df_all = gff3_df.merge(df_two, on=merge_keys, how="left", validate="one_to_one")

#Remove duplicate columns

```

```

        self.data_df = self.__remove_repeated_cols_df__(df_all)

    #drop garbage cols for reduced file size
    self.data_df = self.data_df.drop(columns=['Function_prediction_source','Protein_Lifestyle','Cluster','Subcluster','so'])

    #Filter for only AMR proteins
    if AMR_proteins_Only is True:
        self.data_df = self.data_df[self.data_df["Protein_ID"].isin(phage_AMR_df['Protein_ID'])]

    if save_fname is not None:
        self.savePickle(save_fname,self.data_df)

    return

```

In [2]:

```
'''Example script to organize and merge all of the data for a dataset from PhageScope
For this code to work the following files are required:
```

```
-TemPhD # folder with all of the fasta files
-TemPhD_antimicrobial_resistance_gene_data.tsv
-temphd_phage_annotation_protein_meta_data.tsv
-temphd_phage_meta_data.tsv
-TemPhD.gff3
```

Data can be saved and loaded as .pickle files. There is also an option to export it to .

```
%% Create Data object
# Folder containing all the TSV/GFF and the 'TemPhD' FASTA folder
folder_path = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
db_name = "TemPhD"

data_obj = DataFormat(folder_path, db_name)

%% Generate Data
# Parameters for dataset creation
number_of_AMR_seqs = 10          # number of AMR sequences to include
Number_of_control_seqs = None    # default: 0 non-AMR sequences
high_quality_only = False        # filter only high-quality reads
high_med_quality_only = True     # filter medium & high quality reads
save_fname = None                # None = no automatic save
MR_proteins_Only = True          # only AMR protein sequences

# Call the main function
data_obj.create_data_file(
    number_of_AMR_seqs=number_of_AMR_seqs,
    Number_of_control_seqs=Number_of_control_seqs,
    high_quality_only=high_quality_only,
    high_med_quality_only=high_med_quality_only,
    save_fname=save_fname,
    AMR_proteins_Only=MR_proteins_Only
)

%% Optional: Export merged data to CSV
# data_obj.export_csv("merged_phagescope_data.csv", data_obj.data_df)
```

In [4]:

```
import numpy as np
import pandas as pd
```



```

        'integron', 'cassette'
    ]
s = annot_str.lower()
for kw in keywords:
    if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):
        return True
return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=1.5, label='Viterbi Decoding (Best Path)')
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=1000, step=250, k=4,
    island_radius=3000,           # +/- bp to group ARG-near windows
    min_args_for_island=1,         # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,           # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,        # absolute GC difference threshold
    min_island_length_bp=2000,      # minimum merged island length

```

```

mobile_kw_list=None,
verbose=True
):
"""
Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob
Returns dictionary of results (no CSV writes).
"""

if mobile_kw_list is None:
    mobile_kw_list = [
        'integrase', 'transposase', 'recombinase', 'resolvase',
        'insertion sequence', 'insertion-sequence', 'transposon',
        'integron', 'cassette'
    ]

if verbose:
    print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

if 'DataFormat' not in globals():
    print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code firs
    return

data = DataFormat(folder_path, db_name)

try:
    amr_df = data.__load_tsv__(data.amr_tsv)
    meta_df = data.__load_tsv__(data.meta_tsv)
    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Ph
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_st

if verbose:
    print(f"\n[CHECKPOINT 1] DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries

```

```

features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():
        try:
            s0 = max(0, int(r['Start']) - 1)
            e0 = min(seq_len, int(r['Stop']))
            arg_coords.append((s0, e0))
        except:
            continue
    arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

    # compute mobile positions from annotations
    mobile_positions = []
    for _, prow in my_prots.iterrows():
        annot_fields = []
        for col in ['Product', 'Annotation', 'Function', 'Description', 'product', 'annot']:
            if col in my_prots.columns and pd.notna(prow.get(col, None)):
                annot_fields.append(str(prow.get(col, '')))
        annot_concat = " ".join(annot_fields).strip()
        if is_mobile_element_annotation(annot_concat, mobile_kw_list):
            try:
                ps0 = max(0, int(prow['Start']) - 1)
                pe0 = min(seq_len, int(prow['Stop']))
                mobile_positions.append((ps0, pe0))
            except:
                continue

    # per-phage per-window temporary store
    phage_rows = []

    # sliding windows: collect per-window stats but do NOT finalize label yet
    for i in range(0, max(1, seq_len - window_size + 1), step):
        w_seq = seq[i:i+window_size]
        w_start, w_end = i, min(i + window_size, seq_len)

        # ARG overlaps in window (count)

```

```

arg_overlap_count = 0
for (a0,a1) in arg_coords:
    if a0 < w_end and a1 > w_start:
        arg_overlap_count += 1

# mobile overlaps in window (count)
mobile_overlap_count = 0
for (m0,m1) in mobile_positions:
    if m0 < w_end and m1 > w_start:
        mobile_overlap_count += 1

# near_any_arg: center within island_radius of any ARG midpoint
center = (w_start + w_end)//2
near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

row = {
    'Phage_ID': phage_id,
    'Window_Start': w_start,
    'Window_End': w_end,
    'Local_GC': get_gc(w_seq),
    'kmer': None, # placeholder, we'll expand k-mers later if needed
    'arg_overlap_count': arg_overlap_count,
    'mobile_overlap_count': mobile_overlap_count,
    'near_any_arg': int(near_any_arg),
    # label to be set after merged-range evaluation
    'Label': 0
}
phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].values
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
# (1) region contains >= min_args_for_island ARGs
# (2) region length >= min_island_length_bp
# (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0,m1) in mobile_positions if not (m1 <= rs or m0 >= re))
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
        gc_shift = abs(region_gc - genome_gc)
        region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at least one ARG
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shift >= gc_shift_threshold)
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append((rs, re))

```

```

        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,
            'mobile_count': mobile_count_in_region,
            'gc_shift': gc_shift
        })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End'] <= ar['end'])
        phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list
for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein start/stop
    overlaps = []
    for _, p in my_prots.iterrows():
        try:
            ps0 = max(0, int(p['Start']) - 1)
            pe0 = min(seq_len, int(p['Stop'])))
        except:
            continue
        if ps0 < wend and pe0 > wstart:
            overlaps.append(p)
    overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
    orf_seqs = []
    for _, op in overlapped_prots.iterrows():
        try:
            ps0 = max(0, int(op['Start']) - 1)
            pe0 = min(seq_len, int(op['Stop'])))

```

```

        except:
            continue
        subseq = seq[ps0:pe0]
        if str(op.get('Strand', '+')).strip() in ['-1']:
            subseq = reverse_complement(subseq)
        orf_seqs.append(subseq)
        row_features['Protein_Density'] = len(orf_seqs)

    # simple protein property means if present
    for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
        if c in overlapped_prots.columns and len(overlapped_prots) > 0:
            vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
            row_features[c] = vals.mean() if len(vals) > 0 else 0.0
        else:
            row_features[c] = 0.0

    features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

# Diagnostic printout
if verbose:
    total_with_islands = island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

```

```

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)
    coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:
    print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
    print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
    print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verb
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:

```

```

print("\n==== EVALUATION ON TEST DATASET ===")
print(classification_report(y_test, y_pred_viterbi, target_names=['Background',
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_viterbi)))
print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages))
    for pid in viz_phages:
        mask = (df_test['Phage_ID'] == pid).values
        X_p = X_test[mask]
        windows = df_test[mask]['Window_Start'].values
        y_true_p = df_test[mask]['Label'].values
        y_vit = model.predict(X_p, [len(X_p)])
        try:
            y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
        except:
            y_post = np.zeros(len(X_p))
        visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
else:
    if verbose:
        print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length
df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written.)")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

```

```

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=1000,
    step=250,
    k=4,
    island_radius=3000,
    min_args_for_island=1,
    min_mobile_count=2,
    gc_shift_threshold=0.1,
    min_island_length_bp=2000,
    verbose=True
)

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION
Total Valid Phages: 328
Training Set: 262 | Test Set: 66
-----
Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ====
Phages analyzed: 328. Phages with ≥1 accepted island: 35
Examples (first 10):
- TemPhD_cluster_46242: [{'start': 1250, 'end': 8000, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.032191993632864235}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_39211: [{'start': 37000, 'end': 43750, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.0033462167099317375}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_22243: [{'start': 37000, 'end': 43750, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.003123835229784877}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_56217: [{'start': 33000, 'end': 39750, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.0257332710173091}] | #args=1 mobile_hits=4 genome_GC=0.345
- TemPhD_cluster_22244: [{'start': 37000, 'end': 43750, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.003219978290659453}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_46209: [{'start': 48000, 'end': 54750, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.015212533920434201}] | #args=1 mobile_hits=5 genome_GC=0.354
- TemPhD_cluster_22237: [{'start': 37000, 'end': 43750, 'length': 6750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.003155882916743069}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_54409: [{'start': 27250, 'end': 34000, 'length': 6750, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.000777157353978819}] | #args=1 mobile_hits=4 genome_GC=0.479
- TemPhD_cluster_54261: [{'start': 37500, 'end': 45000, 'length': 7500, 'arg_count': 2, 'mobile_count': 2, 'gc_shift': 0.05513425312915887}] | #args=3 mobile_hits=6 genome_GC=0.359
- TemPhD_cluster_55771: [{'start': 61250, 'end': 70500, 'length': 9250, 'arg_count': 2, 'mobile_count': 3, 'gc_shift': 0.043768527477040686}] | #args=2 mobile_hits=7 genome_GC=0.398

[CHECKPOINT] LASSO selected 202 features.

```

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.7168

[CHECKPOINT 3] MODEL TRAINED


== EVALUATION ON TEST DATASET ==


	precision	recall	f1-score	support
Background	1.00	1.00	1.00	15625
Island	0.81	0.80	0.81	168
accuracy			1.00	15793
macro avg	0.90	0.90	0.90	15793
weighted avg	1.00	1.00	1.00	15793

Confusion Matrix:

```
[[15593  32]
 [ 33 135]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)





Predicted islands after merging (count = 7). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_46242	1750	6250	4500
1	TemPhD_cluster_56564	49000	52750	3750
2	TemPhD_cluster_46209	47750	55250	7500
3	TemPhD_cluster_56212	33000	40250	7250
4	TemPhD_cluster_22242	36750	44250	7500

PIPELINE COMPLETE. Returning results (no CSVs written).

In [5]:

```
import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
```

```

import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====

def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)
    comp = str.maketrans("ACGTacgt", "TGCATgca")
    return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):

```

```

        return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=1.5, label='Viterbi Decoding (Best Path)')
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=1000, step=500, k=4,
    island_radius=3000,                      # +/- bp to group ARG-near windows
    min_args_for_island=1,                     # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,                        # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,                    # absolute GC difference threshold
    min_island_length_bp=2000,                 # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob

```

```

    Returns dictionary of results (no CSV writes).
"""

if mobile_kw_list is None:
    mobile_kw_list = [
        'integrase', 'transposase', 'recombinase', 'resolvase',
        'insertion sequence', 'insertion-sequence', 'transposon',
        'integron', 'cassette'
    ]

if verbose:
    print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

if 'DataFormat' not in globals():
    print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code first")
    return

data = DataFormat(folder_path, db_name)

try:
    amr_df = data.__load_tsv__(data.amr_tsv)
    meta_df = data.__load_tsv__(data.meta_tsv)
    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n[CHECKPOINT 1] DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

```

```

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():
        try:
            s0 = max(0, int(r['Start']) - 1)
            e0 = min(seq_len, int(r['Stop']))
            arg_coords.append((s0, e0))
        except:
            continue
    arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

    # compute mobile positions from annotations
    mobile_positions = []
    for _, prow in my_prots.iterrows():
        annot_fields = []
        for col in ['Product', 'Annotation', 'Function', 'Description', 'product', 'annotation']:
            if col in my_prots.columns and pd.notna(prow.get(col, None)):
                annot_fields.append(str(prow.get(col, '')))
        annot_concat = " ".join(annot_fields).strip()
        if is_mobile_element_annotation(annot_concat, mobile_kw_list):
            try:
                ps0 = max(0, int(prow['Start']) - 1)
                pe0 = min(seq_len, int(prow['Stop']))
                mobile_positions.append((ps0, pe0))
            except:
                continue

    # per-phage per-window temporary store
    phage_rows = []

    # sliding windows: collect per-window stats but do NOT finalize label yet
    for i in range(0, max(1, seq_len - window_size + 1), step):
        w_seq = seq[i:i+window_size]
        w_start, w_end = i, min(i + window_size, seq_len)

        # ARG overlaps in window (count)
        arg_overlap_count = 0
        for (a0,a1) in arg_coords:
            if a0 < w_end and a1 > w_start:
                arg_overlap_count += 1

```

```

# mobile overlaps in window (count)
mobile_overlap_count = 0
for (m0,m1) in mobile_positions:
    if m0 < w_end and m1 > w_start:
        mobile_overlap_count += 1

# near_any_arg: center within island_radius of any ARG midpoint
center = (w_start + w_end)//2
near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

row = {
    'Phage_ID': phage_id,
    'Window_Start': w_start,
    'Window_End': w_end,
    'Local_GC': get_gc(w_seq),
    'kmer': None, # placeholder, we'll expand k-mers later if needed
    'arg_overlap_count': arg_overlap_count,
    'mobile_overlap_count': mobile_overlap_count,
    'near_any_arg': int(near_any_arg),
    # label to be set after merged-range evaluation
    'Label': 0
}
phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].values
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
# (1) region contains >= min_args_for_island ARGs
# (2) region length >= min_island_length_bp
# (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0,m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
        gc_shift = abs(region_gc - genome_gc)
        region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,

```

```

        'mobile_count': mobile_count_in_region,
        'gc_shift': gc_shift
    })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End']
phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list
for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1']:
        subseq = reverse_complement(subseq)

```

```

        orf_seqs.append(subseq)
        row_features['Protein_Density'] = len(orf_seqs)

        # simple protein property means if present
        for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
            if c in overlapped_prots.columns and len(overlapped_prots) > 0:
                vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
                row_features[c] = vals.mean() if len(vals) > 0 else 0.0
            else:
                row_features[c] = 0.0

        features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

# Diagnostic printout
if verbose:
    total_with_islands = island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)

```

```

coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:
    print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
    print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
    print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verb
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n== EVALUATION ON TEST DATASET ==")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background',
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

```

```

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages))
        for pid in viz_phages:
            mask = (df_test['Phage_ID'] == pid).values
            X_p = X_test[mask]
            windows = df_test[mask]['Window_Start'].values
            y_true_p = df_test[mask]['Label'].values
            y_vit = model.predict(X_p, [len(X_p)])
            try:
                y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
            except:
                y_post = np.zeros(len(X_p))
            visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
    else:
        if verbose:
            print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length
df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written.)")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=1000,

```

```

step=500,
k=4,
island_radius=3000,
min_args_for_island=1,
min_mobile_count=2,
gc_shift_threshold=0.1,
min_island_length_bp=2000,
verbose=True
)

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION
Total Valid Phages: 328
Training Set: 262 | Test Set: 66
-----
Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ====
Phages analyzed: 328. Phages with ≥1 accepted island: 14
Examples (first 10):
- TemPhD_cluster_46242: [{'start': 1500, 'end': 8000, 'length': 6500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.03287575431662493}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_54409: [{'start': 27500, 'end': 34000, 'length': 6500, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.0003168597400383111}] | #args=1 mobile_hits=4 genome_GC=0.479
- TemPhD_cluster_54261: [{'start': 37500, 'end': 45000, 'length': 7500, 'arg_count': 2, 'mobile_count': 2, 'gc_shift': 0.05513425312915887}] | #args=3 mobile_hits=6 genome_GC=0.359
- TemPhD_cluster_55771: [{'start': 61500, 'end': 70500, 'length': 9000, 'arg_count': 2, 'mobile_count': 3, 'gc_shift': 0.04482858753710073}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_45788: [{'start': 61500, 'end': 70500, 'length': 9000, 'arg_count': 2, 'mobile_count': 3, 'gc_shift': 0.04458295402889206}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_55772: [{'start': 61500, 'end': 70500, 'length': 9000, 'arg_count': 2, 'mobile_count': 3, 'gc_shift': 0.04490285028582558}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_56167: [{'start': 7000, 'end': 13500, 'length': 6500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.024757012265705203}] | #args=2 mobile_hits=4 genome_GC=0.356
- TemPhD_cluster_45961: [{'start': 11500, 'end': 18000, 'length': 6500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.01686148347571015}] | #args=1 mobile_hits=4 genome_GC=0.345
- TemPhD_cluster_58001: [{'start': 46500, 'end': 53000, 'length': 6500, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.01593907074472445}] | #args=3 mobile_hits=3 genome_GC=0.399
- TemPhD_cluster_55462: [{'start': 46000, 'end': 52500, 'length': 6500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.027468087234893968}] | #args=2 mobile_hits=4 genome_GC=0.299

[CHECKPOINT] LASSO selected 144 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.7563
-----
[CHECKPOINT 3] MODEL TRAINED

```




== EVALUATION ON TEST DATASET ==


	precision	recall	f1-score	support
Background	1.00	0.99	0.99	7879
Island	0.20	0.53	0.29	36
accuracy			0.99	7915
macro avg	0.60	0.76	0.64	7915
weighted avg	0.99	0.99	0.99	7915

Confusion Matrix:

```
[[7802  77]
 [ 17  19]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)





Predicted islands after merging (count = 7). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_46242	2000	6000	4000
1	TemPhD_cluster_56564	49000	52500	3500
2	TemPhD_cluster_46209	48000	55500	7500
3	TemPhD_cluster_56212	33000	40500	7500
4	TemPhD_cluster_22242	37000	45000	8000

PIPELINE COMPLETE. Returning results (no CSVs written).

In [6]:

```

import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====
def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)

```

```

comp = str.maketrans("ACGTacgt", "TGCATgca")
return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):
            return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=2)
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=500, step=250, k=4,
    island_radius=3000,           # +/- bp to group ARG-near windows
    min_args_for_island=1,         # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,           # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,        # absolute GC difference threshold
    min_island_length_bp=2000,     # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob
    Returns dictionary of results (no CSV writes).
    """

    if mobile_kw_list is None:
        mobile_kw_list = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]

    if verbose:
        print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

    if 'DataFormat' not in globals():
        print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code firs"
        return

    data = DataFormat(folder_path, db_name)

    try:
        amr_df = data.__load_tsv__(data.amr_tsv)
        meta_df = data.__load_tsv__(data.meta_tsv)
    
```

```

    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n{CHECKPOINT_1} DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():

```

```

try:
    s0 = max(0, int(r['Start']) - 1)
    e0 = min(seq_len, int(r['Stop']))
    arg_coords.append((s0, e0))
except:
    continue
arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

# compute mobile positions from annotations
mobile_positions = []
for _, prow in my_prots.iterrows():
    annot_fields = []
    for col in ['Product','Annotation','Function','Description','product','annotation']:
        if col in my_prots.columns and pd.notna(prow.get(col, None)):
            annot_fields.append(str(prow.get(col, '')))
    annot_concat = " ".join(annot_fields).strip()
    if is_mobile_element_annotation(annot_concat, mobile_kw_list):
        try:
            ps0 = max(0, int(prow['Start']) - 1)
            pe0 = min(seq_len, int(prow['Stop']))
            mobile_positions.append((ps0, pe0))
        except:
            continue

# per-phage per-window temporary store
phage_rows = []

# sliding windows: collect per-window stats but do NOT finalize label yet
for i in range(0, max(1, seq_len - window_size + 1), step):
    w_seq = seq[i:i+window_size]
    w_start, w_end = i, min(i + window_size, seq_len)

    # ARG overlaps in window (count)
    arg_overlap_count = 0
    for (a0,a1) in arg_coords:
        if a0 < w_end and a1 > w_start:
            arg_overlap_count += 1

    # mobile overlaps in window (count)
    mobile_overlap_count = 0
    for (m0,m1) in mobile_positions:
        if m0 < w_end and m1 > w_start:
            mobile_overlap_count += 1

    # near_any_arg: center within island_radius of any ARG midpoint
    center = (w_start + w_end)//2
    near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

    row = {
        'Phage_ID': phage_id,
        'Window_Start': w_start,
        'Window_End': w_end,
        'Local_GC': get_gc(w_seq),
        'kmer': None, # placeholder, we'll expand k-mers later if needed
        'arg_overlap_count': arg_overlap_count,
        'mobile_overlap_count': mobile_overlap_count,
        'near_any_arg': int(near_any_arg),
        # label to be set after merged-range evaluation
        'Label': 0
    }

```

```

        }

phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].value_counts()
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
#   (1) region contains >= min_args_for_island ARGs
#   (2) region length >= min_island_length_bp
#   (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0, m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
    gc_shift = abs(region_gc - genome_gc)
    region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,
            'mobile_count': mobile_count_in_region,
            'gc_shift': gc_shift
        })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End'] <= ar['end'])
        phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list

```

```

for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1', '-1']:
        subseq = reverse_complement(subseq)
    orf_seqs.append(subseq)
row_features['Protein_Density'] = len(orf_seqs)

# simple protein property means if present
for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
    if c in overlapped_prots.columns and len(overlapped_prots) > 0:
        vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
        row_features[c] = vals.mean() if len(vals) > 0 else 0.0
    else:
        row_features[c] = 0.0

features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

```

```

# Diagnostic printout
if verbose:
    total_with_islands = island_df[island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df[island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)
    coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:

```

```

print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verbose)
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n==== EVALUATION ON TEST DATASET ===")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background', 'Island']))
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages)))
    for pid in viz_phages:
        mask = (df_test['Phage_ID'] == pid).values
        X_p = X_test[mask]
        windows = df_test[mask]['Window_Start'].values
        y_true_p = df_test[mask]['Label'].values
        y_vit = model.predict(X_p, [len(X_p)])
        try:
            y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
        except:
            y_post = np.zeros(len(X_p))
        visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
else:
    if verbose:
        print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length

```

```

df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written).")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=500,
    step=250,
    k=4,
    island_radius=3000,
    min_args_for_island=1,
    min_mobile_count=2,
    gc_shift_threshold=0.1,
    min_island_length_bp=2000,
    verbose=True
)

```

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION

Total Valid Phages: 328

Training Set: 262 | Test Set: 66

Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===

Phages analyzed: 328. Phages with ≥1 accepted island: 10

Examples (first 10):

```
- TemPhD_cluster_46242: [{"start": 1500, "end": 7750, "length": 6250, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.035694215855086475}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_54409: [{"start": 27500, "end": 33750, "length": 6250, "arg_count": 1, "mobile_count": 3, "gc_shift": 0.0013816017984232265}] | #args=1 mobile_hits=4 genome_GC=0.479
- TemPhD_cluster_54261: [{"start": 37750, "end": 44750, "length": 7000, "arg_count": 2, "mobile_count": 2, "gc_shift": 0.05368663408153984}] | #args=3 mobile_hits=6 genome_GC=0.359
- TemPhD_cluster_55771: [{"start": 61500, "end": 70250, "length": 8750, "arg_count": 2, "mobile_count": 3, "gc_shift": 0.04640636531487852}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_45788: [{"start": 61500, "end": 70250, "length": 8750, "arg_count": 2, "mobile_count": 3, "gc_shift": 0.046836922282860294}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_55772: [{"start": 61500, "end": 70250, "length": 8750, "arg_count": 2, "mobile_count": 3, "gc_shift": 0.0465949137778891}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_56167: [{"start": 7250, "end": 13500, "length": 6250, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.022855473804166748}] | #args=2 mobile_hits=4 genome_GC=0.356
- TemPhD_cluster_58001: [{"start": 46750, "end": 53000, "length": 6250, "arg_count": 1, "mobile_count": 3, "gc_shift": 0.014603686129339821}] | #args=3 mobile_hits=3 genome_GC=0.399
- TemPhD_cluster_55462: [{"start": 46250, "end": 52500, "length": 6250, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.030748087234893973}] | #args=2 mobile_hits=4 genome_GC=0.299
- TemPhD_cluster_54237: [{"start": 61750, "end": 68000, "length": 6250, "arg_count": 1, "mobile_count": 3, "gc_shift": 0.0006337077927310864}] | #args=1 mobile_hits=4 genome_GC=0.479
```

[CHECKPOINT] LASSO selected 106 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.7747

[CHECKPOINT 3] MODEL TRAINED


==== EVALUATION ON TEST DATASET ====

	precision	recall	f1-score	support
Background	1.00	0.99	0.99	15877
Island	0.04	0.17	0.06	48
accuracy			0.98	15925
macro avg	0.52	0.58	0.52	15925
weighted avg	0.99	0.98	0.99	15925

Confusion Matrix:

```
[[15658  219]
 [  40    8]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)



Predicted islands after merging (count = 7). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_56564		49000	3500
1	TemPhD_cluster_46209		48250	4500
2	TemPhD_cluster_39990		10250	2000
3	TemPhD_cluster_56212		33250	4750
4	TemPhD_cluster_22242		37250	4750

PIPELINE COMPLETE. Returning results (no CSVs written).

In [10]:

```
import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
```

```

import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====

def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)
    comp = str.maketrans("ACGTacgt", "TGCATgca")
    return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):

```

```

        return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=1.5, label='Viterbi Decoding (Best Path)')
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=1000, step=100, k=4,
    island_radius=3000,                      # +/- bp to group ARG-near windows
    min_args_for_island=1,                     # require >= this many ARGs in candidate region (we k)
    min_mobile_count=2,                        # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,                    # absolute GC difference threshold
    min_island_length_bp=2000,                 # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob

```

```

    Returns dictionary of results (no CSV writes).
"""

if mobile_kw_list is None:
    mobile_kw_list = [
        'integrase', 'transposase', 'recombinase', 'resolvase',
        'insertion sequence', 'insertion-sequence', 'transposon',
        'integron', 'cassette'
    ]

if verbose:
    print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

if 'DataFormat' not in globals():
    print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code first")
    return

data = DataFormat(folder_path, db_name)

try:
    amr_df = data.__load_tsv__(data.amr_tsv)
    meta_df = data.__load_tsv__(data.meta_tsv)
    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n[CHECKPOINT 1] DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

```

```

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():
        try:
            s0 = max(0, int(r['Start']) - 1)
            e0 = min(seq_len, int(r['Stop']))
            arg_coords.append((s0, e0))
        except:
            continue
    arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

    # compute mobile positions from annotations
    mobile_positions = []
    for _, prow in my_prots.iterrows():
        annot_fields = []
        for col in ['Product', 'Annotation', 'Function', 'Description', 'product', 'annotation']:
            if col in my_prots.columns and pd.notna(prow.get(col, None)):
                annot_fields.append(str(prow.get(col, '')))
        annot_concat = " ".join(annot_fields).strip()
        if is_mobile_element_annotation(annot_concat, mobile_kw_list):
            try:
                ps0 = max(0, int(prow['Start']) - 1)
                pe0 = min(seq_len, int(prow['Stop']))
                mobile_positions.append((ps0, pe0))
            except:
                continue

    # per-phage per-window temporary store
    phage_rows = []

    # sliding windows: collect per-window stats but do NOT finalize label yet
    for i in range(0, max(1, seq_len - window_size + 1), step):
        w_seq = seq[i:i+window_size]
        w_start, w_end = i, min(i + window_size, seq_len)

        # ARG overlaps in window (count)
        arg_overlap_count = 0
        for (a0,a1) in arg_coords:
            if a0 < w_end and a1 > w_start:
                arg_overlap_count += 1

```

```

# mobile overlaps in window (count)
mobile_overlap_count = 0
for (m0,m1) in mobile_positions:
    if m0 < w_end and m1 > w_start:
        mobile_overlap_count += 1

# near_any_arg: center within island_radius of any ARG midpoint
center = (w_start + w_end)//2
near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

row = {
    'Phage_ID': phage_id,
    'Window_Start': w_start,
    'Window_End': w_end,
    'Local_GC': get_gc(w_seq),
    'kmer': None, # placeholder, we'll expand k-mers later if needed
    'arg_overlap_count': arg_overlap_count,
    'mobile_overlap_count': mobile_overlap_count,
    'near_any_arg': int(near_any_arg),
    # label to be set after merged-range evaluation
    'Label': 0
}
phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].values
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
# (1) region contains >= min_args_for_island ARGs
# (2) region length >= min_island_length_bp
# (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0,m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
        gc_shift = abs(region_gc - genome_gc)
        region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,

```

```

        'mobile_count': mobile_count_in_region,
        'gc_shift': gc_shift
    })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End']
phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list
for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1']:
        subseq = reverse_complement(subseq)

```

```

        orf_seqs.append(subseq)
        row_features['Protein_Density'] = len(orf_seqs)

        # simple protein property means if present
        for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
            if c in overlapped_prots.columns and len(overlapped_prots) > 0:
                vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
                row_features[c] = vals.mean() if len(vals) > 0 else 0.0
            else:
                row_features[c] = 0.0

        features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

# Diagnostic printout
if verbose:
    total_with_islands = island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)

```

```

coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:
    print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
    print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
    print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verb
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n== EVALUATION ON TEST DATASET ==")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background',
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

```

```

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages))
        for pid in viz_phages:
            mask = (df_test['Phage_ID'] == pid).values
            X_p = X_test[mask]
            windows = df_test[mask]['Window_Start'].values
            y_true_p = df_test[mask]['Label'].values
            y_vit = model.predict(X_p, [len(X_p)])
            try:
                y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
            except:
                y_post = np.zeros(len(X_p))
            visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
    else:
        if verbose:
            print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length
df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written.)")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=1000,

```

```

step=100,
k=4,
island_radius=3000,
min_args_for_island=1,
min_mobile_count=2,
gc_shift_threshold=0.1,
min_island_length_bp=2000,
verbose=True
)

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION
Total Valid Phages: 328
Training Set: 262 | Test Set: 66
-----
Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ====
Phages analyzed: 328. Phages with ≥1 accepted island: 36
Examples (first 10):
- TemPhD_cluster_46242: [{'start': 1100, 'end': 8000, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.03238523034784008}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_39211: [{'start': 36900, 'end': 43800, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.003694042796888286}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_22243: [{'start': 36900, 'end': 43800, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.0037647370011213988}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_56217: [{'start': 32900, 'end': 39800, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.026367731564813146}] | #args=1 mobile_hits=4 genome_GC=0.345
- TemPhD_cluster_22244: [{'start': 36900, 'end': 43800, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.003860880061995975}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_46209: [{'start': 47800, 'end': 54700, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.017473403485651573}] | #args=1 mobile_hits=5 genome_GC=0.354
- TemPhD_cluster_22237: [{'start': 36900, 'end': 43800, 'length': 6900, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.003796784688079591}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_54409: [{'start': 27100, 'end': 34000, 'length': 6900, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.0022360945520464304}] | #args=1 mobile_hits=4 genome_GC=0.479
- TemPhD_cluster_54261: [{'start': 37500, 'end': 45200, 'length': 7700, 'arg_count': 2, 'mobile_count': 2, 'gc_shift': 0.05412819252309825}] | #args=3 mobile_hits=6 genome_GC=0.359
- TemPhD_cluster_55771: [{'start': 61200, 'end': 70600, 'length': 9400, 'arg_count': 2, 'mobile_count': 3, 'gc_shift': 0.042925514251048735}] | #args=2 mobile_hits=7 genome_GC=0.398

[CHECKPOINT] LASSO selected 199 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.7176
-----
[CHECKPOINT 3] MODEL TRAINED

```


==== EVALUATION ON TEST DATASET ====


	precision	recall	f1-score	support
Background	1.00	1.00	1.00	39018
Island	0.65	0.81	0.72	420
accuracy			0.99	39438
macro avg	0.82	0.90	0.86	39438
weighted avg	0.99	0.99	0.99	39438

Confusion Matrix:

```
[[38835 183]
 [ 79 341]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)





Predicted islands after merging (count = 7). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_46242	1800	6400	4600
1	TemPhD_cluster_56564	49000	52800	3800
2	TemPhD_cluster_46209	47600	55400	7800
3	TemPhD_cluster_56212	32800	40500	7700
4	TemPhD_cluster_22242	36700	44500	7800

PIPELINE COMPLETE. Returning results (no CSVs written).

In [12]:

```
import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====
def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)
```

```

comp = str.maketrans("ACGTacgt", "TGCATgca")
return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):
            return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=2)
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=2000, step=250, k=4,
    island_radius=3000,           # +/- bp to group ARG-near windows
    min_args_for_island=1,         # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,           # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,        # absolute GC difference threshold
    min_island_length_bp=2000,     # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob
    Returns dictionary of results (no CSV writes).
    """

    if mobile_kw_list is None:
        mobile_kw_list = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]

    if verbose:
        print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

    if 'DataFormat' not in globals():
        print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code first")
        return

    data = DataFormat(folder_path, db_name)

    try:
        amr_df = data.__load_tsv__(data.amr_tsv)
        meta_df = data.__load_tsv__(data.meta_tsv)
    
```

```

    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n{CHECKPOINT_1} DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():

```

```

try:
    s0 = max(0, int(r['Start']) - 1)
    e0 = min(seq_len, int(r['Stop']))
    arg_coords.append((s0, e0))
except:
    continue
arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

# compute mobile positions from annotations
mobile_positions = []
for _, prow in my_prots.iterrows():
    annot_fields = []
    for col in ['Product','Annotation','Function','Description','product','annotation']:
        if col in my_prots.columns and pd.notna(prow.get(col, None)):
            annot_fields.append(str(prow.get(col, '')))
    annot_concat = " ".join(annot_fields).strip()
    if is_mobile_element_annotation(annot_concat, mobile_kw_list):
        try:
            ps0 = max(0, int(prow['Start']) - 1)
            pe0 = min(seq_len, int(prow['Stop']))
            mobile_positions.append((ps0, pe0))
        except:
            continue

# per-phage per-window temporary store
phage_rows = []

# sliding windows: collect per-window stats but do NOT finalize label yet
for i in range(0, max(1, seq_len - window_size + 1), step):
    w_seq = seq[i:i+window_size]
    w_start, w_end = i, min(i + window_size, seq_len)

    # ARG overlaps in window (count)
    arg_overlap_count = 0
    for (a0,a1) in arg_coords:
        if a0 < w_end and a1 > w_start:
            arg_overlap_count += 1

    # mobile overlaps in window (count)
    mobile_overlap_count = 0
    for (m0,m1) in mobile_positions:
        if m0 < w_end and m1 > w_start:
            mobile_overlap_count += 1

    # near_any_arg: center within island_radius of any ARG midpoint
    center = (w_start + w_end)//2
    near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

    row = {
        'Phage_ID': phage_id,
        'Window_Start': w_start,
        'Window_End': w_end,
        'Local_GC': get_gc(w_seq),
        'kmer': None, # placeholder, we'll expand k-mers later if needed
        'arg_overlap_count': arg_overlap_count,
        'mobile_overlap_count': mobile_overlap_count,
        'near_any_arg': int(near_any_arg),
        # label to be set after merged-range evaluation
        'Label': 0
    }

```

```

        }

phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].value_counts()
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
#   (1) region contains >= min_args_for_island ARGs
#   (2) region length >= min_island_length_bp
#   (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0, m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
    gc_shift = abs(region_gc - genome_gc)
    region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,
            'mobile_count': mobile_count_in_region,
            'gc_shift': gc_shift
        })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End'] <= ar['end'])
        phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list

```

```

for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1', '-1']:
        subseq = reverse_complement(subseq)
    orf_seqs.append(subseq)
row_features['Protein_Density'] = len(orf_seqs)

# simple protein property means if present
for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
    if c in overlapped_prots.columns and len(overlapped_prots) > 0:
        vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
        row_features[c] = vals.mean() if len(vals) > 0 else 0.0
    else:
        row_features[c] = 0.0

features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

```

```

# Diagnostic printout
if verbose:
    total_with_islands = island_df[island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df[island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)
    coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:

```

```

print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verbose)
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n==== EVALUATION ON TEST DATASET ===")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background', 'Island']))
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages)))
    for pid in viz_phages:
        mask = (df_test['Phage_ID'] == pid).values
        X_p = X_test[mask]
        windows = df_test[mask]['Window_Start'].values
        y_true_p = df_test[mask]['Label'].values
        y_vit = model.predict(X_p, [len(X_p)])
        try:
            y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
        except:
            y_post = np.zeros(len(X_p))
        visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
else:
    if verbose:
        print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length

```

```

df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written).")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=2000,
    step=250,
    k=4,
    island_radius=3000,
    min_args_for_island=1,
    min_mobile_count=2,
    gc_shift_threshold=0.1,
    min_island_length_bp=2000,
    verbose=True
)

```

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION

Total Valid Phages: 328

Training Set: 262 | Test Set: 66

Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===

Phages analyzed: 328. Phages with ≥1 accepted island: 72

Examples (first 10):

```

- TemPhD_cluster_56174: [{'start': 47000, 'end': 54750, 'length': 7750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.038304219167846376}] | #args=1 mobile_hits=4 genome_GC=0.304
- TemPhD_cluster_56171: [{'start': 47000, 'end': 54750, 'length': 7750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.037911393044002484}] | #args=1 mobile_hits=4 genome_GC=0.304
- TemPhD_cluster_46242: [{'start': 750, 'end': 8500, 'length': 7750, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.03183357069379611}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_25241: [{'start': 47000, 'end': 54750, 'length': 7750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.037702255752617375}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_39211: [{'start': 36500, 'end': 44250, 'length': 7750, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.006486001656168305}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_23154: [{'start': 47000, 'end': 54750, 'length': 7750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.038208881605327205}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_22243: [{'start': 36500, 'end': 44250, 'length': 7750, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.0064117683241695644}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_23148: [{'start': 47000, 'end': 54750, 'length': 7750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.03822350533671853}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_56217: [{'start': 32500, 'end': 40250, 'length': 7750, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.028853940073462037}] | #args=1 mobile_hits=4 genome_GC=0.345
- TemPhD_cluster_23152: [{'start': 13750, 'end': 21500, 'length': 7750, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.038725010637585267}] | #args=1 mobile_hits=5 genome_GC=0.304

```

[CHECKPOINT] LASSO selected 183 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.8031

[CHECKPOINT 3] MODEL TRAINED

==== EVALUATION ON TEST DATASET ====

	precision	recall	f1-score	support
Background	1.00	1.00	1.00	15326
Island	0.74	0.79	0.76	203
accuracy			0.99	15529
macro avg	0.87	0.89	0.88	15529
weighted avg	0.99	0.99	0.99	15529

Confusion Matrix:


```

[[15268    58]
 [   42   161]]


```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)


Phage: TemPhD_cluster_56212 - Viterbi Decoding (Best Path)




Posterior Decoding (Classification Uncertainty)




Phage: TemPhD_cluster_22242 - Viterbi Decoding (Best Path)



Posterior Decoding (Classification Uncertainty)





Predicted islands after merging (count = 10). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_46242	1750	7250	5500
1	TemPhD_cluster_56564	49000	53500	4500
2	TemPhD_cluster_46209	46750	56500	9750
3	TemPhD_cluster_56212	32000	41250	9250
4	TemPhD_cluster_45874	46000	55500	9500

PIPELINE COMPLETE. Returning results (no CSVs written).

In [13]:

```

import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====
def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)

```

```

comp = str.maketrans("ACGTacgt", "TGCATgca")
return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):
            return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=2)
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=2000, step=500, k=4,
    island_radius=3000,           # +/- bp to group ARG-near windows
    min_args_for_island=1,         # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,           # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,        # absolute GC difference threshold
    min_island_length_bp=2000,     # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob
    Returns dictionary of results (no CSV writes).
    """

    if mobile_kw_list is None:
        mobile_kw_list = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]

    if verbose:
        print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

    if 'DataFormat' not in globals():
        print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code firs"
        return

    data = DataFormat(folder_path, db_name)

    try:
        amr_df = data.__load_tsv__(data.amr_tsv)
        meta_df = data.__load_tsv__(data.meta_tsv)
    
```

```

    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n{CHECKPOINT_1} DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():

```

```

try:
    s0 = max(0, int(r['Start']) - 1)
    e0 = min(seq_len, int(r['Stop']))
    arg_coords.append((s0, e0))
except:
    continue
arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

# compute mobile positions from annotations
mobile_positions = []
for _, prow in my_prots.iterrows():
    annot_fields = []
    for col in ['Product', 'Annotation', 'Function', 'Description', 'product', 'annot']:
        if col in my_prots.columns and pd.notna(prow.get(col, None)):
            annot_fields.append(str(prow.get(col, '')))
    annot_concat = " ".join(annot_fields).strip()
    if is_mobile_element_annotation(annot_concat, mobile_kw_list):
        try:
            ps0 = max(0, int(prow['Start']) - 1)
            pe0 = min(seq_len, int(prow['Stop']))
            mobile_positions.append((ps0, pe0))
        except:
            continue

# per-phage per-window temporary store
phage_rows = []

# sliding windows: collect per-window stats but do NOT finalize label yet
for i in range(0, max(1, seq_len - window_size + 1), step):
    w_seq = seq[i:i+window_size]
    w_start, w_end = i, min(i + window_size, seq_len)

    # ARG overlaps in window (count)
    arg_overlap_count = 0
    for (a0,a1) in arg_coords:
        if a0 < w_end and a1 > w_start:
            arg_overlap_count += 1

    # mobile overlaps in window (count)
    mobile_overlap_count = 0
    for (m0,m1) in mobile_positions:
        if m0 < w_end and m1 > w_start:
            mobile_overlap_count += 1

    # near_any_arg: center within island_radius of any ARG midpoint
    center = (w_start + w_end)//2
    near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

    row = {
        'Phage_ID': phage_id,
        'Window_Start': w_start,
        'Window_End': w_end,
        'Local_GC': get_gc(w_seq),
        'kmer': None, # placeholder, we'll expand k-mers later if needed
        'arg_overlap_count': arg_overlap_count,
        'mobile_overlap_count': mobile_overlap_count,
        'near_any_arg': int(near_any_arg),
        # label to be set after merged-range evaluation
        'Label': 0
    }

```

```

        }

phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].value_counts()
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
#   (1) region contains >= min_args_for_island ARGs
#   (2) region length >= min_island_length_bp
#   (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0, m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
    gc_shift = abs(region_gc - genome_gc)
    region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,
            'mobile_count': mobile_count_in_region,
            'gc_shift': gc_shift
        })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End'] <= ar['end'])
        phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list

```

```

for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1', '-1']:
        subseq = reverse_complement(subseq)
    orf_seqs.append(subseq)
row_features['Protein_Density'] = len(orf_seqs)

# simple protein property means if present
for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
    if c in overlapped_prots.columns and len(overlapped_prots) > 0:
        vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
        row_features[c] = vals.mean() if len(vals) > 0 else 0.0
    else:
        row_features[c] = 0.0

features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

```

```

# Diagnostic printout
if verbose:
    total_with_islands = island_df[island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df[island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)
    coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:

```

```

print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verbose)
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n==== EVALUATION ON TEST DATASET ===")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background', 'Island']))
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages)))
    for pid in viz_phages:
        mask = (df_test['Phage_ID'] == pid).values
        X_p = X_test[mask]
        windows = df_test[mask]['Window_Start'].values
        y_true_p = df_test[mask]['Label'].values
        y_vit = model.predict(X_p, [len(X_p)])
        try:
            y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
        except:
            y_post = np.zeros(len(X_p))
        visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
else:
    if verbose:
        print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length

```

```

df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written).")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=2000,
    step=500,
    k=4,
    island_radius=3000,
    min_args_for_island=1,
    min_mobile_count=2,
    gc_shift_threshold=0.1,
    min_island_length_bp=2000,
    verbose=True
)

```

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION
Total Valid Phages: 328
Training Set: 262 | Test Set: 66

Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ====
Phages analyzed: 328. Phages with ≥1 accepted island: 72
Examples (first 10):

```
- TemPhD_cluster_56174: [{'start': 47000, 'end': 54500, 'length': 7500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.04011067078074959}] | #args=1 mobile_hits=4 genome_GC=0.304
- TemPhD_cluster_56171: [{'start': 47000, 'end': 54500, 'length': 7500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.03970494143109926}] | #args=1 mobile_hits=4 genome_GC=0.304
- TemPhD_cluster_46242: [{'start': 1000, 'end': 8500, 'length': 7500, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.03161421585508645}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_25241: [{'start': 47000, 'end': 54500, 'length': 7500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.03935386865584323}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_39211: [{'start': 36500, 'end': 44000, 'length': 7500, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.005879550043265069}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_23154: [{'start': 47000, 'end': 54500, 'length': 7500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.04001103214296159}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_22243: [{'start': 36500, 'end': 44000, 'length': 7500, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.005805316711266328}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_23148: [{'start': 47000, 'end': 54500, 'length': 7500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.04002565587435292}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_56217: [{'start': 32500, 'end': 40000, 'length': 7500, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.02825178953582763}] | #args=1 mobile_hits=4 genome_GC=0.345
- TemPhD_cluster_23152: [{'start': 14000, 'end': 21500, 'length': 7500, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.0404110321429616}] | #args=1 mobile_hits=5 genome_GC=0.304
```

[CHECKPOINT] LASSO selected 193 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.7949

[CHECKPOINT 3] MODEL TRAINED

==== EVALUATION ON TEST DATASET ====


	precision	recall	f1-score	support
Background	1.00	1.00	1.00	7681
Island	0.74	0.76	0.75	102
accuracy			0.99	7783
macro avg	0.87	0.88	0.87	7783
weighted avg	0.99	0.99	0.99	7783

Confusion Matrix:


```
[[7653  28]
 [ 24  78]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)


Phage: TemPhD_cluster_22234 - Viterbi Decoding (Best Path)




Posterior Decoding (Classification Uncertainty)




Phage: TemPhD_cluster_56212 - Viterbi Decoding (Best Path)



Posterior Decoding (Classification Uncertainty)





Predicted islands after merging (count = 9). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_46242	2000	6500	4500
1	TemPhD_cluster_56564	49000	53000	4000
2	TemPhD_cluster_46209	47000	56500	9500
3	TemPhD_cluster_56212	32000	41000	9000
4	TemPhD_cluster_57210	48500	50500	2000

PIPELINE COMPLETE. Returning results (no CSVs written).

In [14]:

```

import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====
def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)

```

```

comp = str.maketrans("ACGTacgt", "TGCATgca")
return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):
            return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=2)
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=2500, step=625, k=4,
    island_radius=3000,           # +/- bp to group ARG-near windows
    min_args_for_island=1,         # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,           # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,        # absolute GC difference threshold
    min_island_length_bp=2000,     # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob
    Returns dictionary of results (no CSV writes).
    """

    if mobile_kw_list is None:
        mobile_kw_list = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]

    if verbose:
        print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

    if 'DataFormat' not in globals():
        print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code first")
        return

    data = DataFormat(folder_path, db_name)

    try:
        amr_df = data.__load_tsv__(data.amr_tsv)
        meta_df = data.__load_tsv__(data.meta_tsv)
    
```

```

    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n{CHECKPOINT_1} DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():

```

```

try:
    s0 = max(0, int(r['Start']) - 1)
    e0 = min(seq_len, int(r['Stop']))
    arg_coords.append((s0, e0))
except:
    continue
arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

# compute mobile positions from annotations
mobile_positions = []
for _, prow in my_prots.iterrows():
    annot_fields = []
    for col in ['Product','Annotation','Function','Description','product','annotation']:
        if col in my_prots.columns and pd.notna(prow.get(col, None)):
            annot_fields.append(str(prow.get(col, '')))
    annot_concat = " ".join(annot_fields).strip()
    if is_mobile_element_annotation(annot_concat, mobile_kw_list):
        try:
            ps0 = max(0, int(prow['Start']) - 1)
            pe0 = min(seq_len, int(prow['Stop']))
            mobile_positions.append((ps0, pe0))
        except:
            continue

# per-phage per-window temporary store
phage_rows = []

# sliding windows: collect per-window stats but do NOT finalize label yet
for i in range(0, max(1, seq_len - window_size + 1), step):
    w_seq = seq[i:i+window_size]
    w_start, w_end = i, min(i + window_size, seq_len)

    # ARG overlaps in window (count)
    arg_overlap_count = 0
    for (a0,a1) in arg_coords:
        if a0 < w_end and a1 > w_start:
            arg_overlap_count += 1

    # mobile overlaps in window (count)
    mobile_overlap_count = 0
    for (m0,m1) in mobile_positions:
        if m0 < w_end and m1 > w_start:
            mobile_overlap_count += 1

    # near_any_arg: center within island_radius of any ARG midpoint
    center = (w_start + w_end)//2
    near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

    row = {
        'Phage_ID': phage_id,
        'Window_Start': w_start,
        'Window_End': w_end,
        'Local_GC': get_gc(w_seq),
        'kmer': None, # placeholder, we'll expand k-mers later if needed
        'arg_overlap_count': arg_overlap_count,
        'mobile_overlap_count': mobile_overlap_count,
        'near_any_arg': int(near_any_arg),
        # label to be set after merged-range evaluation
        'Label': 0
    }

```

```

        }

phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].value_counts()
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
#   (1) region contains >= min_args_for_island ARGs
#   (2) region length >= min_island_length_bp
#   (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0, m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
    gc_shift = abs(region_gc - genome_gc)
    region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,
            'mobile_count': mobile_count_in_region,
            'gc_shift': gc_shift
        })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End'] <= ar['end'])
        phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list

```

```

for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1', '-1']:
        subseq = reverse_complement(subseq)
    orf_seqs.append(subseq)
row_features['Protein_Density'] = len(orf_seqs)

# simple protein property means if present
for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
    if c in overlapped_prots.columns and len(overlapped_prots) > 0:
        vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
        row_features[c] = vals.mean() if len(vals) > 0 else 0.0
    else:
        row_features[c] = 0.0

features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

```

```

# Diagnostic printout
if verbose:
    total_with_islands = island_df[island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df[island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)
    coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:

```

```

print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verbose)
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n==== EVALUATION ON TEST DATASET ===")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background', 'Island']))
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages)))
    for pid in viz_phages:
        mask = (df_test['Phage_ID'] == pid).values
        X_p = X_test[mask]
        windows = df_test[mask]['Window_Start'].values
        y_true_p = df_test[mask]['Label'].values
        y_vit = model.predict(X_p, [len(X_p)])
        try:
            y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
        except:
            y_post = np.zeros(len(X_p))
        visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
else:
    if verbose:
        print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length

```

```

df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written).")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=2500,
    step=625,
    k=4,
    island_radius=3000,
    min_args_for_island=1,
    min_mobile_count=2,
    gc_shift_threshold=0.1,
    min_island_length_bp=2000,
    verbose=True
)

```

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION

Total Valid Phages: 328

Training Set: 262 | Test Set: 66

Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===

Phages analyzed: 328. Phages with ≥1 accepted island: 75

Examples (first 10):

```
- TemPhD_cluster_56174: [{"start": 46875, "end": 55000, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.03531067078074962}] | #args=1 mobile_hits=4 genome_GC=0.304
- TemPhD_cluster_56907: [{"start": 45000, "end": 53125, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.06216314412184076}] | #args=1 mobile_hits=5 genome_GC=0.484
- TemPhD_cluster_56171: [{"start": 46875, "end": 55000, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.03628955681571466}] | #args=1 mobile_hits=4 genome_GC=0.304
- TemPhD_cluster_46242: [{"start": 625, "end": 8750, "length": 8125, "arg_count": 1, "mobile_count": 3, "gc_shift": 0.03170652354739417}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_25241: [{"start": 46875, "end": 55000, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.035600022501997064}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_34918: [{"start": 6875, "end": 15000, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.0701265220590987}] | #args=1 mobile_hits=4 genome_GC=0.492
- TemPhD_cluster_39211: [{"start": 36250, "end": 44375, "length": 8125, "arg_count": 1, "mobile_count": 3, "gc_shift": 0.009294934658649723}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_23154: [{"start": 46875, "end": 55000, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.03534436547629494}] | #args=1 mobile_hits=5 genome_GC=0.304
- TemPhD_cluster_22243: [{"start": 36250, "end": 44375, "length": 8125, "arg_count": 1, "mobile_count": 3, "gc_shift": 0.009220701326650982}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_23148: [{"start": 46875, "end": 55000, "length": 8125, "arg_count": 1, "mobile_count": 2, "gc_shift": 0.035358989207686264}] | #args=1 mobile_hits=5 genome_GC=0.304
```

[CHECKPOINT] LASSO selected 187 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.8214

[CHECKPOINT 3] MODEL TRAINED

==== EVALUATION ON TEST DATASET ====


	precision	recall	f1-score	support
Background	1.00	0.99	0.99	6090
Island	0.64	0.73	0.68	92
accuracy			0.99	6182
macro avg	0.82	0.86	0.84	6182
weighted avg	0.99	0.99	0.99	6182

Confusion Matrix:


```
[[6053  37]
 [ 25  67]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)

Phage: TemPhD_cluster_56212 - Viterbi Decoding (Best Path)




Posterior Decoding (Classification Uncertainty)




Phage: TemPhD_cluster_56408 - Viterbi Decoding (Best Path)



Posterior Decoding (Classification Uncertainty)





Predicted islands after merging (count = 16). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_56907	46875	50625	3750
1	TemPhD_cluster_46242	1250	6875	5625
2	TemPhD_cluster_56564	49375	53125	3750
3	TemPhD_cluster_26985	9375	12500	3125
4	TemPhD_cluster_59073	10000	13125	3125

PIPELINE COMPLETE. Returning results (no CSVs written).

In [15]:

```

import numpy as np
import pandas as pd
from collections import Counter
import itertools
import warnings
import matplotlib.pyplot as plt
import random
import re
import os

# Sklearn & HMM Imports
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LassoCV
from hmmlearn import hmm

warnings.filterwarnings('ignore')

# =====
# 1. HELPER FUNCTIONS
# =====
def reverse_complement(seq):
    if seq is None:
        return None
    seq = str(seq)

```

```

comp = str.maketrans("ACGTacgt", "TGCATgca")
return seq.translate(comp)[::-1]

def get_kmers(sequence, k=4):
    if not sequence:
        return {}
    sequence = sequence.upper().replace('N', '')
    if len(sequence) < k:
        return {}
    bases = ['A', 'C', 'G', 'T']
    all_kmers = [''.join(p) for p in itertools.product(bases, repeat=k)]
    counts = Counter([sequence[i:i+k] for i in range(len(sequence)-k+1)])
    total = sum(counts.values())
    if total == 0:
        return {kmer: 0 for kmer in all_kmers}
    return {kmer: counts.get(kmer, 0)/total for kmer in all_kmers}

def get_gc(sequence):
    if not sequence:
        return 0.0
    s = sequence.upper()
    return (s.count('G') + s.count('C')) / len(s)

def is_mobile_element_annotation(annot_str, keywords=None):
    """
    Conservative mobile-element annotation checker using word-boundary regex.
    """
    if not isinstance(annot_str, str) or annot_str.strip() == "":
        return False
    if keywords is None:
        keywords = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]
    s = annot_str.lower()
    for kw in keywords:
        if re.search(r'\b' + re.escape(kw.lower()) + r'\b', s):
            return True
    return False

def visualize_decoding(phage_id, windows, y_true, y_pred_viterbi, y_prob_posterior):
    fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(12, 6), sharex=True)
    ax1.fill_between(windows, 0, y_true, color='green', alpha=0.3, step='mid', label='Ground Truth')
    ax1.step(windows, y_pred_viterbi, where='mid', color='red', linestyle='--', linewidth=2)
    ax1.set_ylabel("State (0/1)")
    ax1.set_title(f"Phage: {phage_id} - Viterbi Decoding (Best Path)")
    ax1.legend(loc='upper right')
    ax1.grid(True, alpha=0.3)

    ax2.fill_between(windows, 0, y_true, color='green', alpha=0.3, label='Ground Truth')
    ax2.plot(windows, y_prob_posterior, linewidth=1.5, label='Posterior Probability')
    ax2.axhline(y=0.5, color='gray', linestyle=':', label='Decision Threshold')
    ax2.set_ylabel("Prob(Island)")
    ax2.set_xlabel("Genome Position (bp)")
    ax2.set_title("Posterior Decoding (Classification Uncertainty)")
    ax2.legend(loc='upper right')
    ax2.grid(True, alpha=0.3)

```

```

plt.tight_layout()
plt.show()

def merge_windows_to_ranges(positions, window_size, step):
    """Merge contiguous/touching window start positions into ranges (start, end)."""
    if len(positions) == 0:
        return []
    pos = np.sort(np.array(positions))
    ranges = []
    start = pos[0]
    prev = pos[0]
    for p in pos[1:]:
        if p <= prev + step:
            prev = p
            continue
        else:
            ranges.append((int(start), int(prev + window_size)))
            start = p
            prev = p
    ranges.append((int(start), int(prev + window_size)))
    return ranges

# =====
# 2. MAIN PROGRAM (strict island definition + HMM)
# =====

def run_pipeline(
    folder_path, db_name="TemPhD",
    window_size=900, step=300, k=3,
    island_radius=3000,           # +/- bp to group ARG-near windows
    min_args_for_island=1,         # require >= this many ARGs in candidate region (we k
    min_mobile_count=2,           # require >= this many mobile annotations in region 0
    gc_shift_threshold=0.1,        # absolute GC difference threshold
    min_island_length_bp=2000,     # minimum merged island length
    mobile_kw_list=None,
    verbose=True
):
    """
    Detect ARG-enriched genomic islands: merged contiguous regions with >=1 ARG AND (mob
    Returns dictionary of results (no CSV writes).
    """

    if mobile_kw_list is None:
        mobile_kw_list = [
            'integrase', 'transposase', 'recombinase', 'resolvase',
            'insertion sequence', 'insertion-sequence', 'transposon',
            'integron', 'cassette'
        ]

    if verbose:
        print("--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-")

    if 'DataFormat' not in globals():
        print("CRITICAL ERROR: 'DataFormat' class not found. Please run Part 1 code firs"
        return

    data = DataFormat(folder_path, db_name)

    try:
        amr_df = data.__load_tsv__(data.amr_tsv)
        meta_df = data.__load_tsv__(data.meta_tsv)
    
```

```

    prot_df = data.__load_tsv__(data.protein_tsv)
except Exception as e:
    print(f"Error loading data: {e}")
    return

# normalize strand
if 'Strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['Strand'].astype(str).str.strip()
elif 'strand' in prot_df.columns:
    prot_df['Strand'] = prot_df['strand'].astype(str).str.strip()
else:
    prot_df['Strand'] = '+'
    if verbose:
        print("WARNING: No 'Strand' column found – assuming '+' for all proteins.")

# prepare phage lists
hq_phages = meta_df[meta_df['Completeness'].isin(['High-quality', 'Complete'])]['Phage_id']
arg_phages = amr_df['Phage_id'].unique()
valid_phages = list(set(hq_phages) & set(arg_phages))

if len(valid_phages) == 0:
    print("No valid phages found after filtering completeness & ARG presence.")
    return

train_phages, test_phages = train_test_split(valid_phages, test_size=0.20, random_state=42)

if verbose:
    print(f"\n{CHECKPOINT_1} DATASET PREPARATION")
    print(f"Total Valid Phages: {len(valid_phages)}")
    print(f"Training Set: {len(train_phages)} | Test Set: {len(test_phages)}")
    print("-----")

# collect global features + island summaries
features_list = []
island_summary = []

if verbose:
    print(f"Extracting per-window features from {len(valid_phages)} phages...")

for idx, phage_id in enumerate(valid_phages):
    df_seq = data.__load_fasta__(phage_id)
    if df_seq is None or df_seq.empty:
        continue
    seq = df_seq.iloc[0]['sequence']
    seq_len = len(seq)
    genome_gc = get_gc(seq)

    my_args = amr_df[amr_df['Phage_id'] == phage_id].copy()
    my_prots = prot_df[prot_df['Phage_ID'] == phage_id].copy()

    # arg protein ids
    if 'Protein_id' in my_args.columns:
        arg_prot_ids = set(my_args['Protein_id'].astype(str))
    else:
        arg_prot_ids = set()

    # build coordinates for ARG-associated proteins (if present)
    arg_coords = []
    for _, r in my_prots[my_prots['Protein_ID'].astype(str).isin(arg_prot_ids)].iterrows():

```

```

try:
    s0 = max(0, int(r['Start']) - 1)
    e0 = min(seq_len, int(r['Stop']))
    arg_coords.append((s0, e0))
except:
    continue
arg_midpoints = [(s+e)//2 for (s,e) in arg_coords]

# compute mobile positions from annotations
mobile_positions = []
for _, prow in my_prots.iterrows():
    annot_fields = []
    for col in ['Product', 'Annotation', 'Function', 'Description', 'product', 'annot']:
        if col in my_prots.columns and pd.notna(prow.get(col, None)):
            annot_fields.append(str(prow.get(col, '')))
    annot_concat = " ".join(annot_fields).strip()
    if is_mobile_element_annotation(annot_concat, mobile_kw_list):
        try:
            ps0 = max(0, int(prow['Start']) - 1)
            pe0 = min(seq_len, int(prow['Stop']))
            mobile_positions.append((ps0, pe0))
        except:
            continue

# per-phage per-window temporary store
phage_rows = []

# sliding windows: collect per-window stats but do NOT finalize label yet
for i in range(0, max(1, seq_len - window_size + 1), step):
    w_seq = seq[i:i+window_size]
    w_start, w_end = i, min(i + window_size, seq_len)

    # ARG overlaps in window (count)
    arg_overlap_count = 0
    for (a0,a1) in arg_coords:
        if a0 < w_end and a1 > w_start:
            arg_overlap_count += 1

    # mobile overlaps in window (count)
    mobile_overlap_count = 0
    for (m0,m1) in mobile_positions:
        if m0 < w_end and m1 > w_start:
            mobile_overlap_count += 1

    # near_any_arg: center within island_radius of any ARG midpoint
    center = (w_start + w_end)//2
    near_any_arg = any(abs(mp - center) <= island_radius for mp in arg_midpoints)

    row = {
        'Phage_ID': phage_id,
        'Window_Start': w_start,
        'Window_End': w_end,
        'Local_GC': get_gc(w_seq),
        'kmer': None, # placeholder, we'll expand k-mers later if needed
        'arg_overlap_count': arg_overlap_count,
        'mobile_overlap_count': mobile_overlap_count,
        'near_any_arg': int(near_any_arg),
        # label to be set after merged-range evaluation
        'Label': 0
    }
    phage_rows.append(row)

```

```

        }

phage_rows.append(row)

# convert to df for this phage and find candidate windows near any ARG
phage_df = pd.DataFrame(phage_rows)
candidate_positions = phage_df[phage_df['near_any_arg'] == 1]['Window_Start'].value_counts()
merged_candidate_ranges = merge_windows_to_ranges(candidate_positions, window_size)

# Evaluate each merged candidate range and accept only those that satisfy final
#   (1) region contains >= min_args_for_island ARGs
#   (2) region length >= min_island_length_bp
#   (3) and (mobile_count >= min_mobile_count OR GC shift >= gc_shift_threshold)
accepted_ranges = []
for (rs, re) in merged_candidate_ranges:
    # count ARGs whose midpoints fall in [rs, re]
    arg_count_in_region = sum(1 for mp in arg_midpoints if (mp >= rs and mp < re))
    # count mobile hits overlapping region
    mobile_count_in_region = sum(1 for (m0, m1) in mobile_positions if not (m1 <=
    # compute GC shift for the full region (extract sequence safely)
    try:
        region_seq = seq[rs:re]
        region_gc = get_gc(region_seq)
    except:
        region_gc = genome_gc
    gc_shift = abs(region_gc - genome_gc)
    region_len = re - rs

    # acceptance criteria
    meets_arg = arg_count_in_region >= max(1, min_args_for_island) # ensure at
    meets_mobile_or_gc = (mobile_count_in_region >= min_mobile_count) or (gc_shi
    meets_length = region_len >= min_island_length_bp

    if meets_arg and meets_mobile_or_gc and meets_length:
        accepted_ranges.append({
            'start': rs,
            'end': re,
            'length': region_len,
            'arg_count': arg_count_in_region,
            'mobile_count': mobile_count_in_region,
            'gc_shift': gc_shift
        })

# set labels for windows that fall inside accepted ranges
if len(accepted_ranges) > 0:
    for ar in accepted_ranges:
        mask = (phage_df['Window_Start'] >= ar['start']) & (phage_df['Window_End'] <= ar['end'])
        phage_df.loc[mask, 'Label'] = 1

# record island summary info
island_summary.append({
    'Phage_ID': phage_id,
    'num_accepted_islands': len(accepted_ranges),
    'accepted_islands': accepted_ranges,
    'num_args': len(arg_midpoints),
    'num_mobile_positions': len(mobile_positions),
    'genome_gc': genome_gc
})

# expand k-mers and other features now and append to global list

```

```

for _, prow in phage_df.iterrows():
    wstart = int(prow['Window_Start'])
    wend = int(prow['Window_End'])
    wseq = seq[wstart:wend]
    row_features = {
        'Phage_ID': phage_id,
        'Window_Start': wstart,
        'Label': int(prow['Label']),
        'Local_GC': prow['Local_GC'],
        'Arg_Count_Window': prow['arg_overlap_count'],
        'Mobile_Count_Window': prow['mobile_overlap_count'],
        'Protein_Density': 0 # we'll fill protein density below
    }
    # add kmer features (k up to k param)
    kmer_feats = get_kmers(wseq, k=k)
    row_features.update(kmer_feats)

    # compute overlapped proteins to approximate protein density and protein sta
overlaps = []
for _, p in my_prots.iterrows():
    try:
        ps0 = max(0, int(p['Start']) - 1)
        pe0 = min(seq_len, int(p['Stop']))
    except:
        continue
    if ps0 < wend and pe0 > wstart:
        overlaps.append(p)
overlapped_prots = pd.DataFrame(overlaps) if len(overlaps) else pd.DataFrame()
orf_seqs = []
for _, op in overlapped_prots.iterrows():
    try:
        ps0 = max(0, int(op['Start']) - 1)
        pe0 = min(seq_len, int(op['Stop']))
    except:
        continue
    subseq = seq[ps0:pe0]
    if str(op.get('Strand', '+')).strip() in ['-1', '-1']:
        subseq = reverse_complement(subseq)
    orf_seqs.append(subseq)
row_features['Protein_Density'] = len(orf_seqs)

# simple protein property means if present
for c in ['Aromaticity', 'Instability_index', 'Isoelectric_point']:
    if c in overlapped_prots.columns and len(overlapped_prots) > 0:
        vals = pd.to_numeric(overlapped_prots[c], errors='coerce').dropna()
        row_features[c] = vals.mean() if len(vals) > 0 else 0.0
    else:
        row_features[c] = 0.0

features_list.append(row_features)

# create features dataframe
df = pd.DataFrame(features_list)
if df.empty:
    print("No features extracted. Exiting.")
    return

# create island summary dataframe for diagnostics
island_df = pd.DataFrame(island_summary)

```

```

# Diagnostic printout
if verbose:
    total_with_islands = island_df[island_df['num_accepted_islands'] > 0].shape[0]
    print("\n==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ===")
    print(f"Phages analyzed: {len(island_df)}. Phages with ≥1 accepted island: {total_with_islands}")
    if total_with_islands > 0:
        print("Examples (first 10):")
        sample_rows = island_df[island_df['num_accepted_islands'] > 0].head(10)
        for _, r in sample_rows.iterrows():
            print(f"- {r['Phage_ID']}: {r['accepted_islands']} | #args={r['num_args']}")

# Now split into train/test by phage using earlier train_phages/test_phages
df_train = df[df['Phage_ID'].isin(train_phages)]
df_test = df[df['Phage_ID'].isin(test_phages)]

# Preprocess: drop id columns and impute/scale
drop_cols = ['Phage_ID', 'Label', 'Window_Start']
imputer = SimpleImputer(strategy='mean')
scaler = StandardScaler()

Xtrain_raw = df_train.drop(columns=drop_cols)
Xtest_raw = df_test.drop(columns=drop_cols)

Xtrain_imp = imputer.fit_transform(Xtrain_raw)
Xtest_imp = imputer.transform(Xtest_raw)

Xtrain_scaled = scaler.fit_transform(Xtrain_imp)
Xtest_scaled = scaler.transform(Xtest_imp)

# LASSO for feature selection
lasso = LassoCV(cv=5, random_state=42, max_iter=5000)
lasso.fit(Xtrain_scaled, df_train['Label'].values)

coef_mask = np.abs(lasso.coef_) > 1e-8
if coef_mask.sum() < 5:
    idxs = np.argsort(-np.abs(lasso.coef_))[:5]
    coef_mask = np.zeros_like(lasso.coef_, dtype=bool)
    coef_mask[idxs] = True

selected_indices = np.where(coef_mask)[0]
X_train = Xtrain_scaled[:, selected_indices]
X_test = Xtest_scaled[:, selected_indices]

if verbose:
    print(f"\n[CHECKPOINT] LASSO selected {len(selected_indices)} features.")

# PCA
n_components = min(50, X_train.shape[1])
if n_components < 1:
    print("Not enough features selected for PCA. Exiting.")
    return
pca = PCA(n_components=n_components)
X_train = pca.fit_transform(X_train)
y_train = df_train['Label'].values
X_test = pca.transform(X_test)
y_test = df_test['Label'].values

if verbose:

```

```

print(f"\n[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE")
print(f"PCA Variance Explained: {np.sum(pca.explained_variance_ratio_):.4f}")
print("-----")

# HMM training
X_bg = X_train[y_train == 0]
X_island = X_train[y_train == 1]

if len(X_island) < 5:
    print("CRITICAL WARNING: not enough island examples in training; aborting.")
    return

model = hmm.GaussianHMM(n_components=2, covariance_type="full", init_params="", verbose)
model.startprob_ = np.array([0.99, 0.01])
model.transmat_ = np.array([[0.98, 0.02], [0.05, 0.95]])
model.means_ = np.array([X_bg.mean(axis=0), X_island.mean(axis=0)])
cov_bg = np.cov(X_bg.T) + np.eye(model.means_.shape[1]) * 0.1
cov_island = np.cov(X_island.T) + np.eye(model.means_.shape[1]) * 0.1
model.covars_ = np.array([cov_bg, cov_island])
model.n_features = model.means_.shape[1]

if verbose:
    print(f"\n[CHECKPOINT 3] MODEL TRAINED")

# predict on test
lengths_test = df_test.groupby('Phage_ID').size().values
if X_test.shape[0] == 0:
    print("No test data available after split.")
    return

y_pred_viterbi = model.predict(X_test, lengths_test)

if verbose:
    print("\n==== EVALUATION ON TEST DATASET ===")
    print(classification_report(y_test, y_pred_viterbi, target_names=['Background', 'Island']))
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred_viterbi))
    print("\n[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)")

# Visualize some phages
test_phages_with_island = df_test[df_test['Label'] == 1]['Phage_ID'].unique()
if len(test_phages_with_island) > 0 and verbose:
    viz_phages = random.sample(list(test_phages_with_island), min(3, len(test_phages)))
    for pid in viz_phages:
        mask = (df_test['Phage_ID'] == pid).values
        X_p = X_test[mask]
        windows = df_test[mask]['Window_Start'].values
        y_true_p = df_test[mask]['Label'].values
        y_vit = model.predict(X_p, [len(X_p)])
        try:
            y_post = model.predict_proba(X_p, [len(X_p)])[:, 1]
        except:
            y_post = np.zeros(len(X_p))
        visualize_decoding(pid, windows, y_true_p, y_vit, y_post)
else:
    if verbose:
        print("No ARG-enriched islands found in Test Set to visualize.")

# merge predicted windows into islands (post-HMM) and filter by min length

```

```

df_test = df_test.copy().reset_index(drop=True)
df_test['Predicted'] = y_pred_viterbi.astype(int)
all_pred_islands = []
for pid in df_test['Phage_ID'].unique():
    df_p = df_test[df_test['Phage_ID'] == pid]
    pred_positions = df_p[df_p['Predicted'] == 1]['Window_Start'].values
    merged = merge_windows_to_ranges(pred_positions, window_size=window_size, step=step)
    merged_filtered = [r for r in merged if (r[1] - r[0]) >= min_island_length_bp]
    for (s,e) in merged_filtered:
        all_pred_islands.append({'Phage_ID': pid, 'pred_island_start': s, 'pred_island_end': e})
pred_islands_df = pd.DataFrame(all_pred_islands)

if verbose:
    if pred_islands_df.empty:
        print("\nNo predicted islands after merging + length filtering.")
    else:
        print(f"\nPredicted islands after merging (count = {len(pred_islands_df)}).")
        print(pred_islands_df.head())

if verbose:
    print("\nPIPELINE COMPLETE. Returning results (no CSVs written).")

return {
    'features_df': df,
    'accepted_island_summary': island_df,
    'predicted_islands': pred_islands_df,
    'hmm_model': model,
    'pca': pca,
    'lasso': lasso,
    'scaler': scaler,
    'imputer': imputer
}

# Example run (adjust path)
my_folder = r"C:\Users\subak\OneDrive\Documents\Comp Bio"
outputs = run_pipeline(
    my_folder,
    window_size=900,
    step=300,
    k=3,
    island_radius=3000,
    min_args_for_island=1,
    min_mobile_count=2,
    gc_shift_threshold=0.1,
    min_island_length_bp=2000,
    verbose=True
)

```

--- STARTING STRICT ARG-ENRICHED ISLAND DETECTION + HMM PIPELINE (strand-aware) ---

[CHECKPOINT 1] DATASET PREPARATION
Total Valid Phages: 328

Training Set: 262 | Test Set: 66

Extracting per-window features from 328 phages...

==== DIAGNOSTIC: ACCEPTED ISLANDS (post-merge + criteria) ====
Phages analyzed: 328. Phages with ≥1 accepted island: 31
Examples (first 10):

```
- TemPhD_cluster_46242: [{'start': 1200, 'end': 7800, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.03411118555205617}] | #args=2 mobile_hits=4 genome_GC=0.394
- TemPhD_cluster_39211: [{'start': 37200, 'end': 43800, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.0026795500432650887}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_22243: [{'start': 37200, 'end': 43800, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.002605316711266348}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_22244: [{'start': 37200, 'end': 43800, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.002701459772140924}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_46209: [{'start': 48000, 'end': 54600, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.016182230890131155}] | #args=1 mobile_hits=5 genome_GC=0.354
- TemPhD_cluster_22237: [{'start': 37200, 'end': 43800, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.00263736439822454}] | #args=1 mobile_hits=3 genome_GC=0.367
- TemPhD_cluster_54409: [{'start': 27300, 'end': 33900, 'length': 6600, 'arg_count': 1, 'mobile_count': 3, 'gc_shift': 0.0013731169499384133}] | #args=1 mobile_hits=4 genome_GC=0.479
- TemPhD_cluster_54261: [{'start': 37800, 'end': 45000, 'length': 7200, 'arg_count': 2, 'mobile_count': 2, 'gc_shift': 0.0554286975736033}] | #args=3 mobile_hits=6 genome_GC=0.359
- TemPhD_cluster_55771: [{'start': 61500, 'end': 70500, 'length': 9000, 'arg_count': 2, 'mobile_count': 3, 'gc_shift': 0.04482858753710073}] | #args=2 mobile_hits=7 genome_GC=0.398
- TemPhD_cluster_22235: [{'start': 37200, 'end': 43800, 'length': 6600, 'arg_count': 1, 'mobile_count': 2, 'gc_shift': 0.002685435928661828}] | #args=1 mobile_hits=3 genome_GC=0.367
```

[CHECKPOINT] LASSO selected 52 features.

[CHECKPOINT 2] DIMENSIONALITY REDUCTION COMPLETE
PCA Variance Explained: 0.9998

[CHECKPOINT 3] MODEL TRAINED

==== EVALUATION ON TEST DATASET ===


	precision	recall	f1-score	support
Background	1.00	0.99	0.99	13073
Island	0.35	0.81	0.49	120
accuracy			0.98	13193
macro avg	0.67	0.90	0.74	13193
weighted avg	0.99	0.98	0.99	13193

Confusion Matrix:


```
[[12894  179]
 [  23   97]]
```

[CHECKPOINT 4] VISUALIZING POSTERIOR DECODING (Uncertainty Analysis)


Phage: TemPhD_cluster_46242 - Viterbi Decoding (Best Path)




Posterior Decoding (Classification Uncertainty)




Phage: TemPhD_cluster_56568 - Viterbi Decoding (Best Path)



Posterior Decoding (Classification Uncertainty)





Predicted islands after merging (count = 15). Example:

	Phage_ID	pred_island_start	pred_island_end	length
0	TemPhD_cluster_46242	2100	6300	4200
1	TemPhD_cluster_46242	55800	57900	2100
2	TemPhD_cluster_56564	49200	52500	3300
3	TemPhD_cluster_46209	6000	8100	2100
4	TemPhD_cluster_46209	48000	56400	8400

PIPELINE COMPLETE. Returning results (no CSVs written).

In []: